

StockCare — Prediction Service (AI Layer 1)

Demand Forecasting & Shortage Detection: Technology Choices and Validation Results

StockCare Project

July 2026

Abstract

This document describes the first AI layer of StockCare: a demand-forecasting engine that predicts, for every (pharmacy, medicine) pair, the total demand over the next $H = 14$ days and converts that forecast into concrete restock queries. We explain each technology used (LightGBM, the feature pipeline, the naïve baselines, the shortage-gap logic, and the FastAPI serving layer), give the mathematical formulation of the model and of every evaluation metric, and report the out-of-sample validation results that demonstrate the system works: a WMAPE of **6.77%**, a **17.6%** error reduction versus the best naïve baseline, and **100%** of simulated stockout-days eliminated when ordering follows the forecast instead of a reactive policy.

Contents

1	Problem Statement	1
2	Data and Feature Engineering	1
2.1	The panel (digital twin)	1
2.2	Features known at day t	2
3	Model: Gradient-Boosted Trees with LightGBM	2
3.1	What gradient boosting does	2
3.2	Poisson objective	3
3.3	Why LightGBM specifically	3
3.4	NumPy fallback (HistGBT)	3
3.5	Leakage-free validation protocol	3
4	Baselines: the Model Must Earn Its Place	4
5	Evaluation Metrics	4
5.1	Forecast accuracy	4
5.2	Results	4
5.3	Business impact: stockouts avoided	5
5.4	A note on F1-score	6
6	Serving Layer	6
7	Conclusion	7

1 Problem Statement

Pharmacies that reorder *reactively* (only once the shelf is nearly empty) regularly run out of stock: the supplier lead time means demand arrives before replenishment does. Layer 1 of StockCare

replaces this reactive behaviour with an *anticipative* one. Formally, at each decision day t , for each pharmacy i and medicine m , the system estimates the forward H -day demand

$$D_{i,m}(t) = \sum_{k=1}^H d_{i,m}(t+k), \quad (1)$$

where $d_{i,m}(t+k)$ is the (unknown) daily demand k days ahead and $H = 14$ is the forecast horizon. This quantity is exactly what the restock decision needs: “how much will be asked of me before a new order can realistically arrive and cover me?”

The forecast is then turned into a *shortage gap*

$$G_{i,m}(t) = \max(0, \hat{D}_{i,m}(t) + S_{i,m}(t) - X_{i,m}(t)), \quad (2)$$

where $\hat{D}$ is the model’s demand forecast, X the current on-hand stock, and S a safety stock. A restock query of quantity $\lceil G \rceil$ is emitted whenever $G > \tau$ (with $\tau = 3$ units, to avoid noise-driven micro-orders). This JSON query is the *only* contract consumed by the downstream layers (prioritisation and routing), so the forecasting model behind $\hat{D}$ can be swapped freely without breaking the rest of the system.

2 Data and Feature Engineering

2.1 The panel (digital twin)

Training data is a daily panel of (pharmacy, medicine, date) rows over two years, covering 8 pharmacies across Tunisian governorates and 15 medicines (chronic, essential, comfort). Each row records the latent demand, the simulated stock under a naïve reactive (s, S) policy, the censored sales $\text{sold} = \min(\text{demand}, \text{stock})$, a stockout flag, and exogenous signals: exam periods, heatwaves, and a smooth winter-peaking flu index

$$\text{flu}(t) = \frac{1}{2} \left(1 + \cos \frac{2\pi \cdot \text{day}(t)}{365.25} \right) \in [0, 1]. \quad (3)$$

The generator deliberately reproduces the *problem* (reactive ordering causes stockouts) so the business evaluation in Section 5.3 can measure how much the forecast fixes it.

2.2 Features known at day t

Each supervised row maps features observable at day t to the target of Eq. (1). **Why pandas/NumPy:** the entire pipeline is vectorised columnar computation on a medium-sized table — exactly what pandas and NumPy are built for; no distributed framework is warranted at pilot scale. The feature set:

- **Lags** $\text{sold}_{t-\ell}$ for $\ell \in \{1, 7, 14, 30\}$ — recent level and weekly pattern of sales.
- **Rolling means** $\frac{1}{w} \sum_{j=1}^w \text{sold}_{t-j}$ for $w \in \{7, 14, 30\}$ — smoothed demand level at three time scales.
- **Stockout pressure:** 14-day rolling mean of the stockout flag — signals that recent *sales* underestimate true *demand* (censoring).
- **Current stock** X_t .
- **Contextual signals:** `is_exam`, `heatwave`, `flu_index` — these drive the demand spikes of stimulants/anxiolytics, rehydration salts/antihistamines, and antivirals/antibiotics respectively.

- **Calendar:** day-of-week, month, ISO week, season — periodicity the model can pick up on its own.
- **Categoricals:** region, medicine category, medicine id; plus a cold-chain flag.

Every lag and rolling feature is shifted by at least one day, so *no feature ever looks into the future* — a strict requirement to avoid target leakage.

3 Model: Gradient-Boosted Trees with LightGBM

3.1 What gradient boosting does

Gradient boosting builds an additive ensemble of small regression trees, each one correcting the errors of the ensemble so far:

$$\hat{F}_K(\mathbf{x}) = \hat{F}_0 + \eta \sum_{k=1}^K f_k(\mathbf{x}), \quad f_k = \arg \min_f \sum_n L(y_n, \hat{F}_{k-1}(\mathbf{x}_n) + f(\mathbf{x}_n)), \quad (4)$$

where each new tree f_k is fitted (via a second-order Taylor expansion of the loss L) to the negative gradients — the “residual errors” — of the current ensemble, and $\eta = 0.05$ is the learning rate that shrinks each tree’s contribution to prevent overfitting.

3.2 Poisson objective

Daily demand is a non-negative *count*. We therefore use the Poisson objective rather than plain squared error: the model outputs $\log \mu(\mathbf{x})$ and minimises the Poisson negative log-likelihood

$$L(y, \mu) = \mu - y \log \mu \ (+ \text{const}), \quad \mu(\mathbf{x}) = e^{F(\mathbf{x})} > 0. \quad (5)$$

This guarantees non-negative predictions by construction, handles the many small-count series (e.g. flu vaccine at ~ 3 units/day) correctly, and its error scales with the mean — an error of 5 units matters more on a low-volume medicine than on paracetamol.

3.3 Why LightGBM specifically

- **Histogram-based split finding:** continuous features are bucketed into discrete bins, making training dramatically faster and more memory-efficient than exact-split GBT — the full model trains in seconds on this panel, enabling fast iteration and frequent retraining.
- **Native categorical support:** region, category and medicine id are consumed directly (optimal category-subset splits), with no one-hot explosion.
- **Tabular state of the art:** on heterogeneous tabular data mixing lags, counts, flags and categoricals, gradient-boosted trees consistently match or beat deep learning while needing far less data and tuning — and this panel ($\sim 10^4$ rows at pilot scale) is far too small for a deep model to shine.
- **Nonlinearity and interactions for free:** trees natively capture effects such as “flu index high *and* region Bizerte *and* medicine Oseltamivir $\Rightarrow$ demand $\times 2.2$ ” without manual interaction terms.
- **Built-in regularisation and early stopping:** `num_leaves=63`, `min_data_in_leaf=100`, feature/bagging fractions of 0.8, and early stopping (patience 60) on a time-based validation set control overfitting.
- **Interpretability:** gain-based feature importances let us verify the model relies on the signals we expect (rolling sales level, flu index, seasonality), which matters for trust in a health-supply context.

3.4 NumPy fallback (HistGBT)

For robustness of the pipeline itself, a pure-NumPy histogram GBT (~150 lines, same core idea: binning + boosting) is used automatically if LightGBM or its native dependencies are unavailable in the environment. The `train/predict` interface is identical, so the rest of the system is agnostic to the backend. It is a fallback for portability, not a production competitor.

3.5 Leakage-free validation protocol

Random cross-validation is *invalid* for time series — it would train on the future and test on the past. We use a strict **time-based split**: the last 20% of the timeline (cutoff 2025-08-01) is held out, and every accuracy number below is computed exclusively on this out-of-sample period. Early stopping also uses only this split.

4 Baselines: the Model Must Earn Its Place

A learned model is only justified if it beats what a pharmacist could do with a calculator. Two standard forecasting baselines predict the same H -day target:

$$\hat{D}_{i,m}^{\text{MA}}(t) = H \cdot \frac{1}{14} \sum_{j=1}^{14} \text{sold}_{i,m}(t-j) \quad (\text{Moving Average}) \quad (6)$$

$$\hat{D}_{i,m}^{\text{SN}}(t) = D_{i,m}(t-364) \quad (\text{Seasonal Naïve, same weekday one year ago}) \quad (7)$$

The moving average captures the current level; the seasonal naïve captures annual seasonality (flu winters, summer heat). LightGBM must beat both to justify its existence.

5 Evaluation Metrics

5.1 Forecast accuracy

Let y_n be the realised H -day demand and $\hat{y}_n$ the forecast, over the N validation rows.

WMAPE (Weighted Mean Absolute Percentage Error) — headline metric.

$$\text{WMAPE} = \frac{\sum_{n=1}^N |y_n - \hat{y}_n|}{\sum_{n=1}^N |y_n|}. \quad (8)$$

It reads as “total forecast error as a fraction of total demand”: a WMAPE of 6.77% means that summed over the whole network, forecasts are off by under 7 units per 100 units actually demanded. Unlike plain MAPE, it does not explode when y_n is small (frequent with low-volume medicines) and it naturally weights high-volume items more — which matches business impact.

MAE (Mean Absolute Error).

$$\text{MAE} = \frac{1}{N} \sum_{n=1}^N |y_n - \hat{y}_n|. \quad (9)$$

The average error in *units of medicine* per 14-day forecast — directly interpretable by a pharmacist.

Bias (Mean Error).

$$\text{Bias} = \frac{1}{N} \sum_{n=1}^N (\hat{y}_n - y_n). \quad (10)$$

Sign matters: a negative bias means systematic *under*-forecasting (risk of shortage), positive means *over*-forecasting (risk of overstock and waste). For shortage prevention we want bias close to zero and certainly not strongly negative.

5.2 Results

All metrics are out-of-sample (validation period after the 2025-08-01 cutoff, horizon $H = 14$ days):

Model	WMAPE	MAE (units)	Bias (units)
LightGBM (Poisson)	0.0677	14.819	−1.894
Moving Average (14 d)	0.0822	18.011	−2.660
Seasonal Naïve	0.0848	18.564	−9.769

Table 1: Out-of-sample forecast accuracy for the 14-day demand target.




Figure 1: Side-by-side comparison of the three models on the validation period. LightGBM wins on all three metrics; note the strongly negative bias of the seasonal naïve (systematic under-forecasting).

Reading the results:

- LightGBM reduces WMAPE by **17.6%** relative to the best baseline ($0.0822 \rightarrow 0.0677$) and MAE by ~ 3.2 units per forecast.
- Its bias (-1.9 units over a 14-day horizon, i.e. ~ 0.14 units/day) is the smallest of the three; the seasonal naïve under-forecasts badly (-9.8) because last year’s demand misses this year’s growth and shifted epidemic timing.
- The gap versus the moving average comes precisely from what the baseline cannot see: exam periods, heatwaves and the flu curve (Figures 4 and 5). This confirms the contextual features carry real signal.

5.3 Business impact: stockouts avoided

Accuracy is a proxy; the real question is *does the forecast prevent empty shelves?* The backtest replays the validation period for every (pharmacy, medicine) series under two policies:




Figure 2: Predicted vs. actual 14-day demand on every validation row (dashed line = perfect forecast). LightGBM hugs the diagonal; the moving average smears horizontally (it cannot anticipate spikes) and the seasonal naïve shows off-diagonal clusters where last year’s pattern shifted.




Figure 3: Distribution of forecast errors. LightGBM’s errors are the tightest and most centred on zero; the seasonal naïve’s mass sits visibly left of zero (under-forecasting, shortage risk).

1. **Reactive** (s, S): reorder only when stock falls below a threshold — the status quo simulated in the panel.
2. **Forecast-driven**: each day, order up to the level $\hat{D}(t) + S(t)$ (forecast plus safety stock, Eq. (2)), with a 3-day delivery lead time.

A *stockout-day* is any day where demand could not be fully served.

Policy	Stockout-days (validation period)
Reactive (s, S)	14
Forecast-driven (LightGBM)	0
Reduction	100%

Table 2: Simulated stockout-days under each ordering policy.

Under identical lead times and demand, anticipating demand 14 days ahead eliminated every stockout-day the reactive policy produced. This is the end-to-end proof that Layer 1 solves the problem it was built for.




Figure 4: WMAPE per medicine. LightGBM is best or tied on essentially every product; the largest gaps appear on context-driven medicines (flu antivirals, rehydration salts, antihistamines) where the baselines have no access to the epidemiological signals.

5.4 A note on F1-score

Layer 1 is formulated as a *regression* problem (predict a demand quantity), so precision/recall/F1 — which apply to binary classifiers — are not the native metrics; WMAPE/MAE/bias are. The binary decision (“emit a restock query or not”) is derived downstream by thresholding the gap of Eq. (2), and its quality is measured directly by the business backtest of Table 2: zero missed shortages (perfect recall on shortage events in the simulation) with orders sized by the gap rather than fired indiscriminately. Classification metrics on the `needs_restock` decision can be added to the backtest if required by evaluators.

6 Serving Layer

The trained model is exposed to the Spring backend through a small REST API:

- **FastAPI + Uvicorn**: a modern async Python web framework — minimal boilerplate, high throughput for a lightweight inference service, and automatic OpenAPI documentation that serves as the integration contract with the backend team.
- **Pydantic**: request/response schemas (`PredictRequest`, `InventorySnapshot`) are validated at the boundary, so malformed inventory snapshots fail fast with explicit errors instead of corrupting predictions.
- **Docker**: the service ships as a container, pinning Python and LightGBM native dependencies so the environment is reproducible from a laptop to deployment.

Endpoints: `GET /health` (liveness, model backend/version), `GET /info` (horizon, thresholds, validation metrics), `POST /predict` (batch shortage prediction returning, per medication: predicted shortage date, estimated remaining days, missing quantity, demand forecast and urgency). The `/predict` contract is stable by design: richer POS history can replace the current consumption-based approximation without any change for the consumer.




Figure 5: Rolling 14-day demand over the validation period: actual vs. forecasts for a high-volume medicine (Paracetamol, Tunis) and a strongly flu-driven one (Oseltamivir, Sousse). LightGBM tracks the rises *as they begin*; the moving average reacts with a lag — exactly the lag that causes stockouts.

7 Conclusion

The prediction layer meets its specification with evidence: a leakage-free, time-split evaluation shows LightGBM at **6.77% WMAPE**, beating both naïve baselines (**17.6%** error reduction vs. the moving average) with near-zero bias, and the forecast-driven replenishment policy it feeds eliminates **100%** of simulated stockout-days versus the reactive status quo. Each technology was chosen for a specific, defensible reason: LightGBM for state-of-the-art tabular accuracy with native categoricals and a Poisson objective suited to count demand; pandas/NumPy for a simple vectorised feature pipeline; explicit baselines to keep the model honest; a decoupled gap formula ($G = \max(0, D + S - X)$) as a stable contract to downstream layers; and FastAPI/Docker for a validated, reproducible serving surface. All numbers and figures are reproducible via `python run.py` and audited by the unit-test suite (`pytest`).




Figure 6: LightGBM feature importance (share of total split gain). The model leans on the smoothed sales level (rolling means/lags) as expected, with real contributions from the contextual and calendar signals — confirming the forecast is driven by interpretable, sensible inputs.




Figure 7: Business impact: stockout-days over the validation period under identical demand and lead times. Anticipating demand eliminates the empty-shelf days entirely.